


Production Agentics: The Business Asset

Lens 3 of 3 — what agents become when the business owns the substrate they run on

By Robert Paddock · idigdata



I have lived through five platform shifts from the operator side.

The Commodore 64 was the first one I could touch. The graphical browser changed what a business could see. I had an iPhone on day one because I knew the phone had stopped being a phone. I ran a Bitcoin miner when most people still thought the whole thing was a strange computer-science rumor. In 2024, I put agentic work into production.

The technology changed every time. The operator question did not.

Whose asset is this?

That question sounds simple until the business is under pressure. Does the company own the data the system trusts, or does the vendor? Does the business own the decision trail, or does it live in a partner's method? Does the workflow belong to the people who run it, or to a project team that leaves after go-live? When a new platform arrives, the

market asks whether the technology is real. Operators have to ask the harder question: whether the business can hold it, absorb it, and still own itself afterward.

Agentics is the largest shift I have seen — the first where judgment-shaped work, not just data or connection or compute, can be delegated into the substrate of the business. It is also the easiest one to misunderstand.

Agentics is not a transformation by itself. It is an accelerant. It compounds on an owned asset and produces expensive theater on a rented one. If the business has governed data, visible decisions, real ownership, and people who can delegate, verify, and own the result, agents can become business assets. If those foundations are missing, agents make the missing foundation move faster.

That is the hope and the risk, both at once.

Deployment Is Not Outcome

The market has already moved past the question of whether agents can be deployed. They can, and they are. Snowflake is positioning itself as a control plane for the "agentic enterprise" and reported \$1.33 billion in product revenue last quarter, up 34 percent year over year. The platforms are real and the spending is real.

The gap is whether deployment becomes an operating outcome.

A company can have a prompt, a model, access to data, and a live customer-facing workflow, and still miss the business outcome. The agent may answer quickly. It may reduce queue pressure. It may look impressive in a launch announcement. But if the customer experience degrades, if the workflow has no intent, if quality is not governed, if escalation is unclear, or if nobody owns the consequence of the decision, the business has not created an asset. It has installed capability without enough operating judgment around it.

The failures are already on the record. Starbucks rolled out an AI inventory agent across roughly eleven thousand North American stores and retired it about nine months later: it miscounted, confused one product for another, and by several accounts got less accurate over time, until workers were recounting what it scanned. The capability was deployed. The outcome was negative, written across eleven thousand locations. And this failure mode is older than the agents — Australia's Robodebt scheme automated debt assessment against flawed logic and bad data long before the current wave, scaling wrongful debts until the courts and a royal commission unwound it. The technology is new. The pattern is not. An agent does not escape it. It accelerates it.

That distinction is the whole article. Enterprise buyers do not buy capability. They buy risk reduction, operating leverage, and decisions they can stand behind. A pilot can show capability. A business asset has to survive Finance, Operations, Compliance, IT, the customer, and the people who will touch it every day.

The Data Problem Wearing An AI Costume

Agentics does not repeal the old failure pattern. It accelerates it.

Every transformation veteran has seen the shape: the system is technically right and operationally wrong. The report is accurate against bad master data. The workflow matches the process map but not the way work actually moves. The new approval path exists, but the real authority still lives in a spreadsheet, a side conversation, or one person's memory.

An agent does not make that better by being faster.

You cannot automate dysfunction. You can only make it move faster.

This is why the first production-agentics question is whether the business can *hold* the agents. By hold, I mean the ability to give agents governed data, a shared version of

reality, clear ownership of what they may touch, and a decision boundary the business can defend. If the customer master is disputed, the agent automates the dispute. If item meaning is hidden inside tribal codes, the agent scales the confusion. If approval limits are informal, the agent guesses where authority lives. If the shadow spreadsheet is the real decision layer, the agent learns from a source the business cannot govern.

The work starts before automation. It starts by externalizing judgment.

Process is the steps. Judgment is the reasoning behind which steps matter when. Which exceptions can be handled routinely? Which require escalation? Which data source wins when two systems disagree? Which answer is allowed to bind the business? Which human owns the call when the agent is uncertain?

Those questions are not philosophical. They are production requirements. They are also change-management requirements one substrate higher than the ERP era. Agentic-era change is not just training people on new screens. It is training the business on new authority, new trust mechanics, and new evidence standards.

Answer them and agentic capability can become decision integrity: decisions the business can stand behind, trace, and defend. Skip them and the company has automated disagreement at scale.

The Maturity Question

Not every company needs the same agentic shape. The practical question is not "Are we doing AI?" It is: where are we, where do we need to be, and what path can we actually absorb?

The first shape is personal productivity: people using general-purpose AI tools for their own work. That can be useful. It requires policy, training, data-leakage discipline, and judgment. It does not change the operating model.


The second shape is team-level agents around files, folders, documents, and bounded workflows. A finance team may use an agent to prepare reconciliation candidates. A legal team may use one to summarize contract changes. An operations team may use one to assemble supplier confirmations. These tools can create real productivity, but they usually sit beside the operating fabric rather than inside it. By operating fabric, I mean the business-owned combination of data, workflows, systems, decisions, governance, and evidence that lets the company run and improve its work.

The third shape is where agents become production components inside that fabric. They act on governed data. They operate inside bounded decision rights. They leave a trail. They escalate consequential calls to humans. They are monitored for drift. They are owned by named people, not by enthusiasm.

That is where the promise lives, and where the preconditions get strict. Production agents are structurally dependent on the data foundation. If the business does not own the data floor, the agent has nothing stable to stand on. If workflows are not visible, the agent cannot know what it is augmenting. If decisions are not recorded, the agent cannot produce a defensible trail. If people cannot verify the output, the capability will not survive contact with the work.

Three shapes of agentic maturity

Where are you, where do you need to be, and what can you actually absorb?



Personal and team productivity are useful. Production agenticity is structurally dependent on the owned substrate — the asset from Article 2.

I am not describing that third shape from the outside. I run my own practice on it: a governed build environment that ships its own work, a business-owned operating asset that runs a real book of business, and an operations layer with a human on every consequential call. Three production-grade systems, carrying real work, that I would rather walk you through than describe. Eating my own cooking is the point — I do not bring a model to a business I have not first trusted to run mine.

This is not a product menu. It is an honesty test. Some companies should stay at personal and team productivity until the substrate catches up. Some can put narrow agents into production because the decision class is bounded and the evidence is strong. Some are ready to make agenticity part of the operating model. The mistake is pretending those are the same engagement because the demo looks similar.

The Cut That Hasn't Healed

There is one agentics risk I do not know how to close neatly, and I do not trust anyone who claims they can.

What happens when agentics is used internally by malicious intent?

Agentics is a force multiplier for bad too.

That is not a policy footnote. It is a cut that has not healed. Most agentics-risk discussions focus on accidental failure: the model was wrong, the prompt was weak, the retrieval was bad, the guardrail missed. Those are real problems. But deliberate misuse is a different axis. A perfectly obedient agent in the hands of a malicious internal actor can move fast and leave a clean-looking trail around a bad purpose. Alignment at the model layer is necessary, but it is not sufficient.

The market is starting to treat this as infrastructure, not opinion. In 2026 Palo Alto Networks acquired Portkey, an AI gateway, specifically to build a control plane for autonomous agents — the framing being that an agent is a highly privileged insider that has to be monitored, routed, and secured like one. The security industry has worked out that you cannot protect what you cannot see.

This is also where the hostile buyer question belongs, because it is the right question: when the agent acts and it is wrong, who owns the failure, and what is the ROI on something non-deterministic?


My answer starts with ownership. The accountability does not live in the model. It lives in the operator-side audit substrate: the decision-rights matrix that says what the agent may and may not decide, the audit trail that records what happened and why, drift detection that surfaces behavior departing from spec, and a human on every consequential call. The AI advisor Neil Sahota put the legal version of this plainly: the first agent that gets a company sued will not be the AI's fault, it will be the company's,

because "the vendor caused it" is not a defense that survives a regulator or a courtroom. The accountability was always the operator's. The substrate is what lets the operator carry it.

An agent you cannot watch is an agent you cannot govern. This is what observability is for, and it is the discipline most agentic programs name last. Observability is the operator-side layer that sits above the substrate and makes the agent's behavior legible — what it read, what it decided, what it touched, where it drifted. It is what turns "a human on the call" from a slogan into a control: the human can observe, validate, and approve because the system was built to be seen. The same security industry now buying agent control planes arrived at this from its own direction — you cannot protect what you cannot see. Watching the agent is not surveillance theater. It is the precondition for letting it act at all.

The Architectural Datum

Hold the substrate. Absorb it into the people. Watch the agent from above.



ANSWER BOTH → DECISION INTEGRITY

Decisions the business can stand behind, trace, and defend. Skip the questions and you automate the disagreement at scale.

An agent you cannot watch is an agent you cannot govern.

Hold the substrate, absorb it into the people, and watch the agent from the datum above — that is what makes capability into decision integrity.

Non-determinism is not governed by pretending the system is deterministic. It is governed by bounding what the agent may decide, preserving the trail of what it did, and assigning ownership before the failure occurs. The point is not to make the system predictable. It is to make the consequences bounded and the evidence complete.

That is the partial answer. It is not a complete moral one. Governance can make malicious use harder, slower, more visible, and more attributable. Segregation of duties can stop one person from turning intent into action without another control point. Decision trails make patterns visible. Drift detection flags behavior that no longer matches the approved envelope. A human on every consequential call keeps the company from pretending an agent should carry moral responsibility. But no serious operator should claim governance makes malicious intent impossible. The risk remains. The discipline is to design the substrate so bad action has to cross named boundaries, leave evidence, and meet human accountability sooner rather than later.

That posture is not fear. It is what lets the hope be honest.

The Workforce Does Not Follow The Org Chart

The easy workforce story says agentics either expands people or displaces them.

That is too clean.

I misjudged who would adapt. The encouragement I gave was the same for everyone, and it was light — a conversation about this new world called agentics, go try Claude, see what it does for your work. The technical people all said yes. Only about half actually adapted. Other people, in other departments, took the same nudge and built genuinely useful apps and real domain expertise on their own — fast enough that we quickly had to put governance around what they were building. The same

encouragement produced grassroots production work in one place and nothing in another, and the org chart, the title, and the technical credential did not predict which.

That is still not solved.

Adding agentic apps into the technology stack and putting governance around them is the part I know how to do. Identifying who in the company is best suited for agentic work is harder, and agentics is new enough that the pattern is still forming. It will remove certain job functions from some people and empower job functions for others. The honest operator has to hold both, and watch for the quiet version of the same problem the shadow spreadsheet always represented: the moment someone's self-built agent becomes load-bearing, the business has a production system it never governed.

This is the *absorb* half of the frame. Hold asks whether the business has the governed substrate. Absorb asks whether the people can delegate, verify, challenge, and own the output. A human on every consequential call is only meaningful if that human understands what they are approving, when to override, and when the agent is producing confidence instead of truth.

The hopeful part is real. When agentics lands well, work feels different. Routine cognitive load shrinks. People spend less time moving data between systems and more time applying judgment. Newer people reach institutional context faster. Experienced operators extend their reach. The company takes on work it could not previously reach without adding proportional headcount.

But the hope is not "everyone gets superpowers." The honest version is that some people turn the tools into leverage quickly, some need structured support, and some work changes enough that the old job shape no longer holds. The operator's responsibility is not to pretend the mapping is obvious. It is to find the people who can absorb the capability, protect the business from misuse, train the judgment layer, and make workforce decisions explicitly rather than letting the tool decide by accident.

The Function Above The Engineer

The market is recognizing that agentic deployment needs people in the field — and it is reaching for a model that already has a name. The Forward Deployed Engineer was Palantir's invention, going back to the early 2010s: send the engineer into the customer's environment because the work could not be understood from a distance. That model is now being scaled into the agentic era by the labs. OpenAI launched a deployment company in May 2026, backed by more than four billion dollars, and acquired the firm Tomoro to bring roughly a hundred and fifty Forward Deployed Engineers in from day one. The same month, Anthropic launched a one-and-a-half-billion-dollar venture with Blackstone, Hellman & Friedman, and Goldman Sachs to deploy Claude across the portfolio companies of those private-equity backers. The signal matters more than the figures: the labs know that model capability is not enough, and they are buying their way into the field.

Forward Deployed Engineers matter. They bring the build closer to the work. They are still only half the deployment-to-outcome gap.

The other half is what I call the Forward Deployed Operator — the FDO is my term, the deliberate level above the FDE. Where the engineer builds the agent, the operator owns the business outcome: the decision class the agent is allowed to touch, the economics, the control boundary, the absorption path through the workforce, and the accountability when something goes wrong. That is not a program manager with a new title. A program manager runs the plan; the FDO owns whether the capability becomes an asset the business can carry after everyone leaves.

That distinction matters because agentic is the accelerant, not the identity. My claim is not that I am the best engineer in the room, or that idigdata owns the tool stack. The claim is that production agentic fails when engineering is not aimed by someone who understands Finance, Operations, Compliance, IT, data ownership, change absorption, and the politics of real work. The whole agentic discourse has split into two rooms —

the people shipping agents and the people trying to govern them — and they are converging on the same conclusion from both directions: the differentiator is no longer the agent. It is the operating model and the governance fabric around it. That fabric is operator work.

This also answers the question a private-equity-backed operator is right to ask: my sponsor is already pushing one of these lab ventures into our operating companies — why an independent operator instead? Because those ventures are distribution channels with a financial incentive to embed one model and one stack. They are the engineer half, sold through the sponsor. The FDO is the half that stays accountable to the business, not the vendor: vendor-agnostic, keeping the company's data and the resulting asset in the company's hands, and aiming whatever engineers show up — theirs, yours, or anyone's. One operator at the apex is not outgunned by a hundred-and-fifty-engineer subsidiary. The engineers are an aimable resource. The accountable apex is the scarce part.

Diagnosers explain. Builders demo. Operators finish.

The finishing function is what buyers should test for. Can the person in front of you name the business decision class the agent is allowed to touch? Can they explain who owns the failure? Can they show how the evidence trail will work before go-live? Can they say what the human validates and what the agent may decide? Can they tell you when not to deploy the agent yet? An engineer can build the agent. An FDO makes sure the business is ready to own what the agent becomes.

Agentics As An Asset

Pilots die in the P&L. Assets live on the balance sheet.

That is not a trick of accounting language. It is a discipline of intent. Article 2 made the mechanics point: a transformation becomes fundable as an asset only when the work creates something the business can defend, run, and carry. Agentics should be held to the same standard. If the agent is a temporary experiment, a vendor demo, a sidecar chatbot, or a departmental tool with no durable evidence trail, it is operating expense. Sometimes that is the right treatment.

And the boundary has to be honest, because an auditor will be. The volatile parts — the constant model swapping, the prompt iteration, the exploratory tuning — are not the asset; they are running cost, and treating them as capital is exactly the move that gets a transformation a credibility problem with its own CFO. What can become an asset is the durable build: the governed data foundation, the decision-rights and audit substrate, the workflows and integrations the business will run on for years. The CFO owns that judgment and the auditor has to agree; the operator's job is to make the evidence clean enough that the judgment is defensible — what was built, why, what decision class it serves, which labor and vendor work created the asset, when it moved into service, and what is training or change-management rather than build. The full mechanics live in Article 2.

Then comes the question a CFO always asks, and should: what is the ROI? I will not hand you a fabricated number. Enterprise agentics — the production beat the industry is leaning into now — is barely two years old. Nobody has a clean multi-year return for it yet, and anyone who shows you a confident five-year ROI is selling you the confidence, not the return.

But the data already runs in both directions. The negative ROI is measured. The first generation — the rush to be first, the corporate moonshots, the glorified chatbots bolted onto ungoverned operations — largely failed, and Starbucks wrote one of those failures across eleven thousand stores. That is what capability without substrate costs, and it lands in the P&L. The positive return is turning now, and it turns precisely where the substrate is real: bounded decision classes, owned data, an evidence trail. The way

you will know your own ROI is not "AI productivity" in the abstract — it is concrete and operator-measurable. For a reconciliation agent it is fewer unreconciled items aging past sixty days, fewer analyst hours per resolved exception, and a cleaner audit trail at close. For an intake agent it is throughput per reviewer and a lower error rate on what reaches the system of record. You measure it against the metric you already own.

The most honest ROI I can give you is my own. Agentics has made one operator into the output of a team — I run a practice, ship production systems, and carry work that used to need more people, because I built it on a substrate I own. That is the shape of the return when it is real. It is also the warning: the multiplier works just as hard on a weak foundation, in the wrong direction.

This is where agentics stops being a technology initiative and starts being a capital discipline. The question is not "Can we expense another pilot?" It is "Are we building a business asset that will keep producing operating leverage after the consultant, the vendor, and the launch announcement are gone?" That question forces ownership to be defined early, forces the agent to attach to real workflows, forces Finance into the room before the story is retrofitted, and forces the evidence to exist while the work is happening. The discipline that makes agentics defensible is the same discipline that makes it valuable.

The Thread Underneath

The thread underneath all five platform shifts is data.

Not data as a warehouse project. Not data as a dashboard. Data as the operating truth of the business: the definitions, ownership, workflows, exceptions, decisions, and evidence that let people and systems act from the same reality.

I want to be careful with the word governance, because it has been hollowed out. In too many rooms it means a committee, an access list, a policy nobody reads — group-think dressed up as control. For data that misses the point. Governing data is not locking fields in a database or approving them by committee. It is making the data mean something the business owns and can stand behind: definitions, provenance, ownership, the shared version of reality that lets a person and an agent act from the same truth. You do not control data into trustworthiness. You dig down until the business can see itself clearly enough that the data is true — and then the agent has something real to stand on.

That is why this practice is called idigdata. The same digging that makes a transformation hold is what makes production agentic possible. The same people foundation from Article 1 and mechanics foundation from Article 2 become the substrate agents either compound or expose.

So the operator question remains the same: whose asset is this? If the answer is the vendor's, the platform's, the partner's, or one heroic employee's private logic, agentic will magnify the dependency. If the answer is the operating company's, agentic can magnify the company.

I am hopeful about this shift because I have seen what happens when the substrate is real. Agents can take routine burden out of the day, force evidence into the decision trail, and give people governed access to knowledge that used to live in one person's memory — letting a mid-market company carry complexity that used to require more headcount, more meetings, and more fragile workarounds.

I am also cautious, because the same force multiplier works in the other direction. Bad data moves faster. Bad intent gets leverage. Bad workforce assumptions get expensive. Bad ownership becomes visible only after the business is already exposed. Holding both is the work.

Hope is earned through proof, not hype. The proof is not that the agents are impressive. The proof is that the business can hold them, absorb them, account for them, govern them, and keep the asset when the outside help leaves.

The agents are ready. The question is whether the business is ready to own what they touch.



Bring the real situation. If you are deciding what to let agents touch, start a conversation at idigdata.com.